

▼

Description of the problem - machine learning libraries and packages experimental setup

```
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.svm import SVC
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.preprocessing import StandardScaler, MinMaxScaler
import seaborn as sns
import time

from zipfile import ZipFile

from sklearn.model_selection import RandomizedSearchCV

# configuring the path of Kaggle.json file
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json

# API to fetch the dataset from Kaggle
!kaggle datasets download -d faresabbasai2022/heart-diseases-prediction-with-streamlit

Downloading heart-diseases-prediction-with-streamlit.zip to /content
 0% 0.00/719k [00:00<?, ?0/s]
100% 719k/719k [00:00<00:00, 130MB/s]

# extracting the compressed Dataset
data = '/content/heart-diseases-prediction-with-streamlit.zip'

with ZipFile(data,'r') as zip:
    zip.extractall()
    print('The dataset is extracted')

    The dataset is extracted

!ls

drive      heart-diseases-prediction-with-streamlit.zip  kaggle.json
heart.csv  heart.jpeg                                   sample_data
```

▼

Choice of dataset Data Mining

```
# Load the dataset
data_path = '/content/heart.csv'
Heart_data = pd.read_csv(data_path)

# Display the first few rows to confirm it's loaded correctly
print(Heart_data.head())

   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  \
0    52    1   0      125   212    0         1      168     0         1.0     2
1    53    1   0      140   203    1         0      155     1         3.1     0
2    70    1   0      145   174    0         1      125     1         2.6     0
3    61    1   0      148   203    0         1      161     0         0.0     2
4    62     0   0      138   294    1         1      106     0         1.9     1

   ca  thal  target
0    2     3       0
1    0     3       0
2    0     3       0
3    1     3       0
4    3     2       0

print(Heart_data.isnull().sum())

age          0
sex          0
cp           0
trestbps     0
chol         0
fbs          0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64

# drop the missing data
Heart_data = Heart_data.dropna()

# the shape after dropping the missing data
Heart_data.shape


(1027, 14)

# Showing the data after Converting categorical values to numeric values
Heart_data.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0

Next steps:

Generate code with Heart\_data

 View recommended plots

```
# Original data frame had 10 columns, we now have 14 columns
Heart_data.shape

(1027, 14)

# Displaying statistical information about the dataset
print(Heart_data.describe())

   age          sex          cp          trestbps          chol  \
count  1027.000000  1027.000000  1027.000000  1027.000000  1027.000000
mean    54.411879     0.695229     0.940604    131.530672    245.764362
std      9.144326     0.460035     1.029476     17.595625     51.817785
min     18.000000     0.000000     0.000000     90.000000    125.000000
25%     48.000000     0.000000     0.000000    120.000000    211.000000
50%     56.000000     1.000000     1.000000    130.000000    240.000000
75%     61.000000     1.000000     2.000000    140.000000    275.000000
max     77.000000     1.000000     3.000000    200.000000    564.000000

   fbs          restecg          thalach          exang          oldpeak  \
count  1027.000000  1027.000000  1027.000000  1027.000000  1027.000000
mean     0.148978     0.528724    148.960078     0.335930     1.069426
std     0.356240     0.527880     23.246693     0.472545     1.174858
min      0.000000     0.000000     70.000000     0.000000     0.000000
25%      0.000000     0.000000     132.000000     0.000000     0.000000
50%      0.000000     1.000000    152.000000     0.000000     0.800000
75%      0.000000     1.000000    166.000000     1.000000     1.800000
max      1.000000     3.000000    202.000000     1.000000     6.200000

   slope          ca          thal          target
count  1027.000000  1027.000000  1027.000000  1027.000000
mean    1.382668     0.755599     2.320351     0.513145
std      0.620171     1.032444     0.625642     0.500071
min      0.000000     0.000000     0.000000     0.000000
25%      1.000000     0.000000     2.000000     0.000000
50%      1.000000     0.000000     2.000000     1.000000
75%      2.000000     1.000000     3.000000     1.000000
max      2.000000     4.000000     3.000000     1.000000

# Information about data types and non-null counts
print(Heart_data.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1027 entries, 0 to 1026
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1027 non-null   int64
1   sex         1027 non-null   int64
2   cp          1027 non-null   int64
3   trestbps    1027 non-null   int64
4   chol        1027 non-null   int64
5   fbs         1027 non-null   int64
6   restecg     1027 non-null   int64
7   thalach     1027 non-null   int64
8   exang       1027 non-null   int64
```


```
9    oldpeak    1027 non-null    float64
10   slope      1027 non-null    int64
11   ca         1027 non-null    int64
12   thal       1027 non-null    int64
13   target     1027 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 112.5 KB
None
```

▼ Correlation Analysis:

To check how each feature correlates with the target variable

```
# Calculate the correlation matrix
corr_matrix = Heart_data.corr()

# Use seaborn to create a heatmap to visualize the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap='coolwarm', cbar=True)
plt.show()
```



▼ Splitting the Data into Training and Testing Sets

```
# 'thalach (Maximum Heart Rate Achieved): Highest heart rate achieved.' is the target variable
X = Heart_data.drop('thalach', axis=1) # Features
y = Heart_data['thalach']              # Target
```

```
# Splitting the dataset
X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

▼ Data Normalization

```
# Normalize features
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Choice of machine learning techniques

Optimization/Parametrization

Models Training

▼ 1- Linear Regression

```
# Train the Linear Regression model
lin_reg = LinearRegression()
lin_reg.fit(X_train, Y_train)
```

```
LinearRegression
LinearRegression()
```

▼ 2-Random Forest Regressor

```
# Random Forest Regressor
from sklearn.ensemble import RandomForestRegressor
rf_reg = RandomForestRegressor(random_state=42)
rf_reg.fit(X_train_scaled, Y_train)
```

```
RandomForestRegressor
RandomForestRegressor(random_state=42)
```

▼ 3- Gradient Boosting Regressor

```
# Gradient Boosting Regressor
from sklearn.ensemble import GradientBoostingRegressor
gb_reg = GradientBoostingRegressor(random_state=42)
gb_reg.fit(X_train_scaled, Y_train)
```

```
GradientBoostingRegressor
GradientBoostingRegressor(random_state=42)
```

Evaluate the performance of the machine learning methods metrics

▼ Evaluate Model

```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Assuming rf_reg and gb_reg are already trained...
```

```
# Predictions
lin_predictions = lin_reg.predict(X_test)
rf_predictions = rf_reg.predict(X_test_scaled)
gb_predictions = gb_reg.predict(X_test_scaled)
```

```
# Calculate the metrics for each model
lin_mse = mean_squared_error(Y_test, lin_predictions)
lin_mae = mean_absolute_error(Y_test, lin_predictions)
lin_r2 = r2_score(Y_test, lin_predictions)
```

```
rf_mse = mean_squared_error(Y_test, rf_predictions)
rf_mae = mean_absolute_error(Y_test, rf_predictions)
rf_r2 = r2_score(Y_test, rf_predictions)
```

```
gb_mse = mean_squared_error(Y_test, gb_predictions)
gb_mae = mean_absolute_error(Y_test, gb_predictions)
gb_r2 = r2_score(Y_test, gb_predictions)
```


```
# Create a dictionary with the model names and their corresponding metrics
model_metrics = {
    'Model': ['Linear Regression', 'Random Forest Regressor', 'Gradient Boosting Regressor'],
    'MSE': [lin_mse, rf_mse, gb_mse],
    'MAE': [lin_mae, rf_mae, gb_mae],
    'R2': [lin_r2, rf_r2, gb_r2]
}
```

```
# Convert the dictionary to a DataFrame
comparison_df = pd.DataFrame(model_metrics)
```


```
# Display the DataFrame
print(comparison_df)
```

	Model	MSE	MAE	R2
0	Linear Regression	283.690883	13.259960	0.462960
1	Random Forest Regressor	18.895370	2.681845	0.964230
2	Gradient Boosting Regressor	115.191848	8.174055	0.781937


```
plt.figure(figsize=(10, 6))
plt.scatter(Y_test, lin_predictions, alpha=0.6)
plt.plot([Y_test.min(), Y_test.max()], [Y_test.min(), Y_test.max()], 'k--', lw=4) # Diagonal line
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs. Predicted Values( Linear Regression)')
plt.show()
```



```
plt.figure(figsize=(10, 6))
plt.scatter(Y_test, rf_predictions, alpha=0.6, color='green') # Using green for differentiation
plt.plot([Y_test.min(), Y_test.max()], [Y_test.min(), Y_test.max()], 'k--', lw=4) # Diagonal line
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs. Predicted Values (Random Forest Regressor)')
plt.show()
```



```
plt.figure(figsize=(10, 6))
plt.scatter(Y_test, gb_predictions, alpha=0.6, color='blue') # Using blue for differentiation
plt.plot([Y_test.min(), Y_test.max()], [Y_test.min(), Y_test.max()], 'k--', lw=4) # Diagonal line
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs. Predicted Values (Gradient Boosting Regressor)')
plt.show()
```



Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.