

Description of the problem - machine learning libraries and packages experimental

setup

```
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.svm import SVM
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.preprocessing import StandardScaler, MinMaxScaler
import seaborn as sns
import time

from zipfile import ZipFile

from sklearn.model_selection import RandomizedSearchCV

# configuring the path of Kaggle.json file
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json

# API to fetch the dataset from Kaggle
!kaggle datasets download -d dileep070/heart-disease-prediction-using-logistic-regression

Downloading heart-disease-prediction-using-logistic-regression.zip to /content
 0% 0.00/58.4k [00:00<?, ?B/s]
100% 58.4k/58.4k [00:00<00:00, 73.3MB/s]

# extracting the compressed Dataset
data = '/content/heart-disease-prediction-using-logistic-regression.zip'

with ZipFile(data,'r') as zip:
    zip.extractall()
    print('The dataset is extracted')

    The dataset is extracted

!ls

framingham.csv  heart-disease-prediction-using-logistic-regression.zip  kaggle.json  sample_data
```

Choice of dataset Data Mining

```
# Load the dataset
data_path = '/content/framingham.csv'
framingham_data = pd.read_csv(data_path)

# Display the first few rows to confirm it's loaded correctly
print(framingham_data.head())

   male  age  education  currentSmoker  cigsPerDay  BPMeds  prevalentStroke  \
0      1   39         4.0             0           0.0      0.0              0.0  \
1      0   46         2.0             0           0.0      0.0              0.0  \
2      1   48         1.0             1          20.0      0.0              0.0  \
3      0   61         3.0             1          30.0      0.0              0.0  \
4      0   46         3.0             1          23.0      0.0              0.0  \

   prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  heartRate  glucose  \
0              0         0      195.0  106.0   70.0   26.97     80.0    77.0  \
1              0         0      250.0  121.0   81.0   28.73     95.0    76.0  \
2              0         0      245.0  127.5   80.0   25.34     75.0    70.0  \
3              1         0      225.0  150.0   95.0   28.58     65.0   103.0  \
4              0         0      285.0  130.0   84.0   23.10     85.0    85.0  \

   TenYearCHD
0              0
1              0
2              0
3              1
4              0

# drop the missing data
Heart_data = framingham_data.dropna()

# the shape after dropping the missing data
framingham_data.shape

(4238, 16)

print(framingham_data.isnull().sum())

male          0
age            0
education     105
currentSmoker  0
cigsPerDay    29
BPMeds        53
prevalentStroke  0
prevalentHyp  0
diabetes       0
totChol       50
sysBP         0
diaBP         0
BMI           19
heartRate     1
glucose       388
TenYearCHD    0
dtype: int64

# drop the missing data
framingham_data = framingham_data.dropna()

# the shape after dropping the missing data
framingham_data.shape

(3656, 16)

print(framingham_data.isnull().sum().sort_values(ascending=False))

male          0
age            0
education     105
currentSmoker  0
cigsPerDay    29
BPMeds        53
prevalentStroke  0
prevalentHyp  0
diabetes       0
totChol       50
sysBP         0
diaBP         0
BMI           19
heartRate     1
glucose       388
TenYearCHD    0
dtype: int64

# Showing the data after Converting categorical values to numeric values
framingham_data.head()
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	TenYearCHD
0	1	39	4.0	0	0.0	0.0	0	0	0	195.0	106.0	70.0	26.97	80.0	77.0	0
1	0	46	2.0	0	0.0	0.0	0	0	0	250.0	121.0	81.0	28.73	95.0	76.0	0
2	1	48	1.0	1	20.0	0.0	0	0	0	245.0	127.5	80.0	25.34	75.0	70.0	0
3	0	61	3.0	1	30.0	0.0	0	1	0	225.0	150.0	95.0	28.58	65.0	103.0	1
4	0	46	3.0	1	23.0	0.0	0	0	0	285.0	130.0	84.0	23.10	85.0	85.0	0

Next steps: [Generate code with framingham_data](#) [View recommended plots](#)

```
# Displaying statistical information about the dataset
print(framingham_data.describe())

   male  age  education  currentSmoker  cigsPerDay  \
count  3656.000000  3656.000000  3656.000000  3656.000000  3656.000000  \
mean    0.443654  49.557440  1.979759  0.489059  9.022155  \
std     0.496883  8.561133  1.022657  0.499949  11.918869  \
min     0.000000  32.000000  1.000000  0.000000  0.000000  \
25%     0.000000  42.000000  1.000000  0.000000  0.000000  \
50%     0.000000  49.000000  2.000000  0.000000  0.000000  \
75%     1.000000  56.000000  3.000000  1.000000  20.000000  \
max      1.000000  70.000000  4.000000  1.000000  70.000000  \

   BPMeds  prevalentStroke  prevalentHyp  diabetes  totChol  \
count  3656.000000  3656.000000  3656.000000  3656.000000  3656.000000  \
mean    0.030361    0.005744    0.311543    0.027079   236.873085  \
std     0.171602    0.075581    0.463187    0.162335   44.096223  \
min     0.000000    0.000000    0.000000    0.000000   113.000000  \
25%     0.000000    0.000000    0.000000    0.000000   206.000000  \
```

50%	0.000000	0.000000	0.000000	0.000000	234.000000
75%	0.000000	0.000000	0.000000	1.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	600.000000

	sysBP	diaBP	BMI	heartRate	glucose	\
count	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000	
mean	132.368025	82.912062	25.784185	75.730580	81.856127	
std	22.092444	11.974825	4.065913	11.982952	23.910128	
min	83.500000	48.000000	15.540000	44.000000	40.000000	
25%	117.000000	75.000000	23.080000	68.000000	71.000000	
50%	128.000000	82.000000	25.380000	75.000000	78.000000	
75%	144.000000	90.000000	28.040000	82.000000	87.000000	
max	295.000000	142.500000	56.800000	143.000000	394.000000	

	TenYearCHD
count	3656.000000
mean	0.152352
std	0.359411
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
# Information about data types and non-null counts
print(framingham_data.info())
```


```
<class 'pandas.core.frame.DataFrame'>
Index: 3656 entries, 0 to 4237
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   male             3656 non-null   int64
1   age              3656 non-null   int64
2   education        3656 non-null   float64
3   currentSmoker    3656 non-null   int64
4   cigsPerDay       3656 non-null   float64
5   BPMeds           3656 non-null   float64
6   prevalentStroke  3656 non-null   int64
7   prevalentHyp     3656 non-null   int64
8   diabetes         3656 non-null   int64
9   totChol          3656 non-null   float64
10  sysBP            3656 non-null   float64
11  diaBP            3656 non-null   float64
12  BMI              3656 non-null   float64
13  heartRate        3656 non-null   float64
14  glucose          3656 non-null   float64
15  TenYearCHD       3656 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 485.6 KB
None
```

▼ Correlation Analysis:

To check how each feature correlates with the target variable

```
# Calculate the correlation matrix
corr_matrix = framingham_data.corr()

# Use seaborn to create a heatmap to visualize the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap='coolwarm', cbar=True)
plt.show()
```



▼ Splitting the Data into Training and Testing Sets

```
# heartRate (Maximum Heart Rate Achieved) is the target variable
X = Heart_data.drop('heartRate', axis=1) # Features
y = Heart_data['heartRate']              # Target
```

```
# Splitting the dataset
X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

▼ Data Normalization

```
# Normalize features
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Choice of machine learning techniques

Optimization/Parametrization

Models Training

▼ 1- Linear Regression

```
# Train the Linear Regression model
lin_reg = LinearRegression()
lin_reg.fit(X_train, Y_train)
```

```
+ LinearRegression
LinearRegression()
```

▼ 2-Random Forest Regressor

```
# Random Forest Regressor
from sklearn.ensemble import RandomForestRegressor
rf_reg = RandomForestRegressor(random_state=42)
rf_reg.fit(X_train_scaled, Y_train)
```

```
+ RandomForestRegressor
RandomForestRegressor(random_state=42)
```

▼ 3- Gradient Boosting Regressor

```
# Gradient Boosting Regressor
from sklearn.ensemble import GradientBoostingRegressor
gb_reg = GradientBoostingRegressor(random_state=42)
gb_reg.fit(X_train_scaled, Y_train)
```

```
+ GradientBoostingRegressor
GradientBoostingRegressor(random_state=42)
```

Evaluate the performance of the machine learning methods metrics

▼ Evaluate Model

```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Assuming rf_reg and gb_reg are already trained...

# Predictions
lin_predictions = lin_reg.predict(X_test)
rf_predictions = rf_reg.predict(X_test_scaled)
gb_predictions = gb_reg.predict(X_test_scaled)

# Calculate the metrics for each model
lin_mse = mean_squared_error(Y_test, lin_predictions)
lin_mae = mean_absolute_error(Y_test, lin_predictions)
lin_r2 = r2_score(Y_test, lin_predictions)

rf_mse = mean_squared_error(Y_test, rf_predictions)
rf_mae = mean_absolute_error(Y_test, rf_predictions)
rf_r2 = r2_score(Y_test, rf_predictions)

gb_mse = mean_squared_error(Y_test, gb_predictions)
gb_mae = mean_absolute_error(Y_test, gb_predictions)
gb_r2 = r2_score(Y_test, gb_predictions)


# Create a dictionary with the model names and their corresponding metrics
model_metrics = {
    'Model': ['Linear Regression', 'Random Forest Regressor', 'Gradient Boosting Regressor'],
    'MSE': [lin_mse, rf_mse, gb_mse],
    'MAE': [lin_mae, rf_mae, gb_mae],
    'R2': [lin_r2, rf_r2, gb_r2]
}

# Convert the dictionary to a DataFrame
comparison_df = pd.DataFrame(model_metrics)

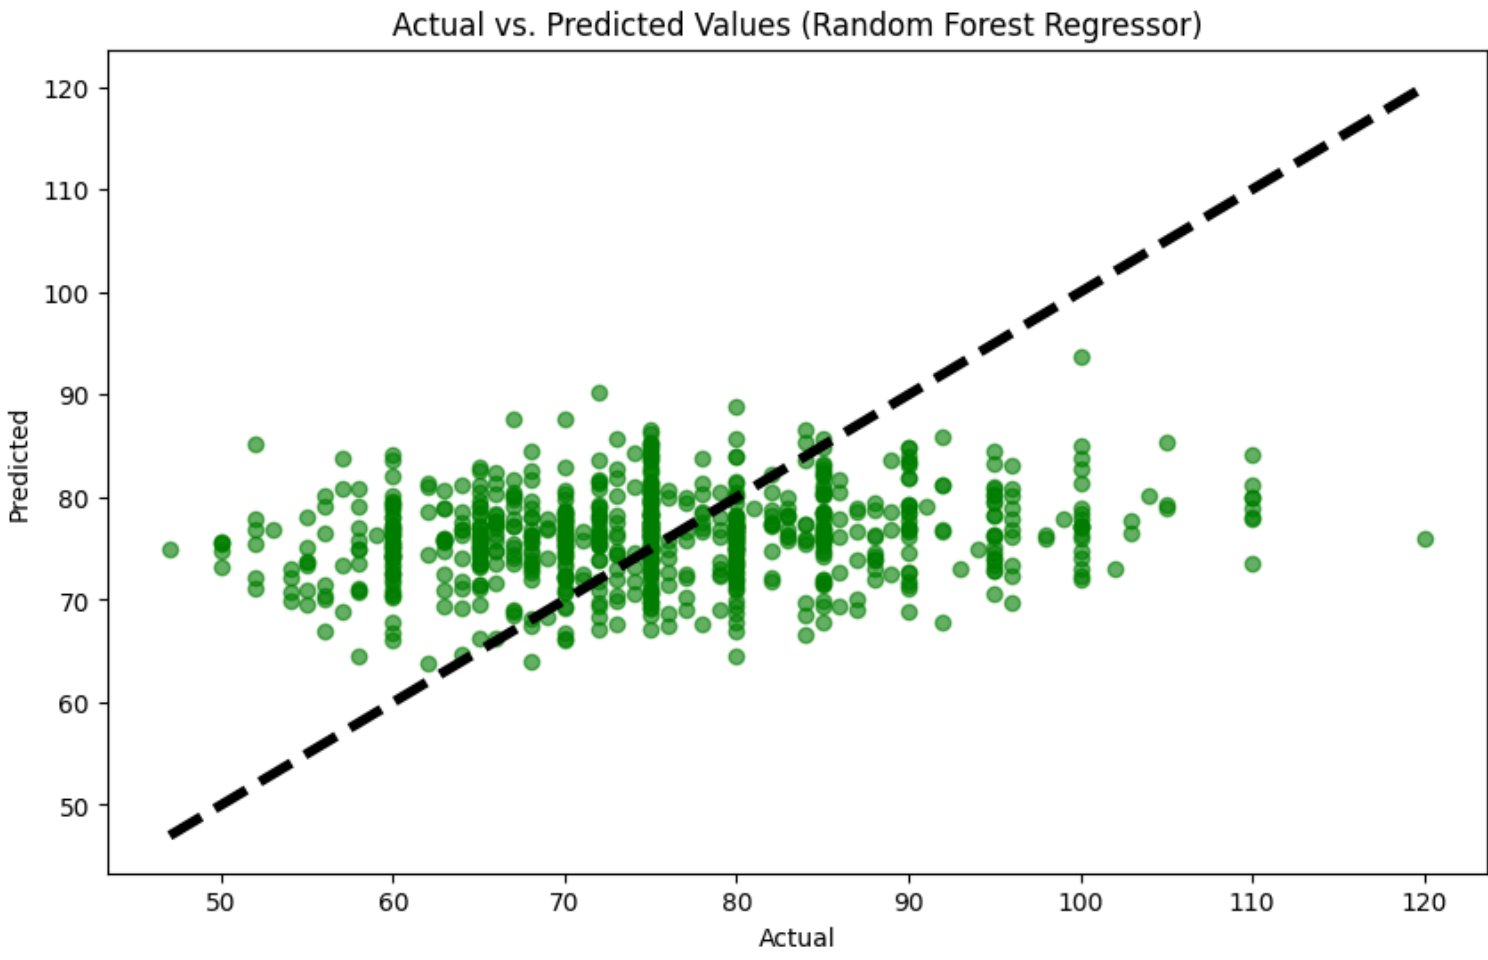
# Display the DataFrame
print(comparison_df)
```

	Model	MSE	MAE	R2
0	Linear Regression	137.609384	9.377027	0.059675
1	Random Forest Regressor	145.445384	9.691954	0.006129
2	Gradient Boosting Regressor	138.934627	9.465155	0.050619

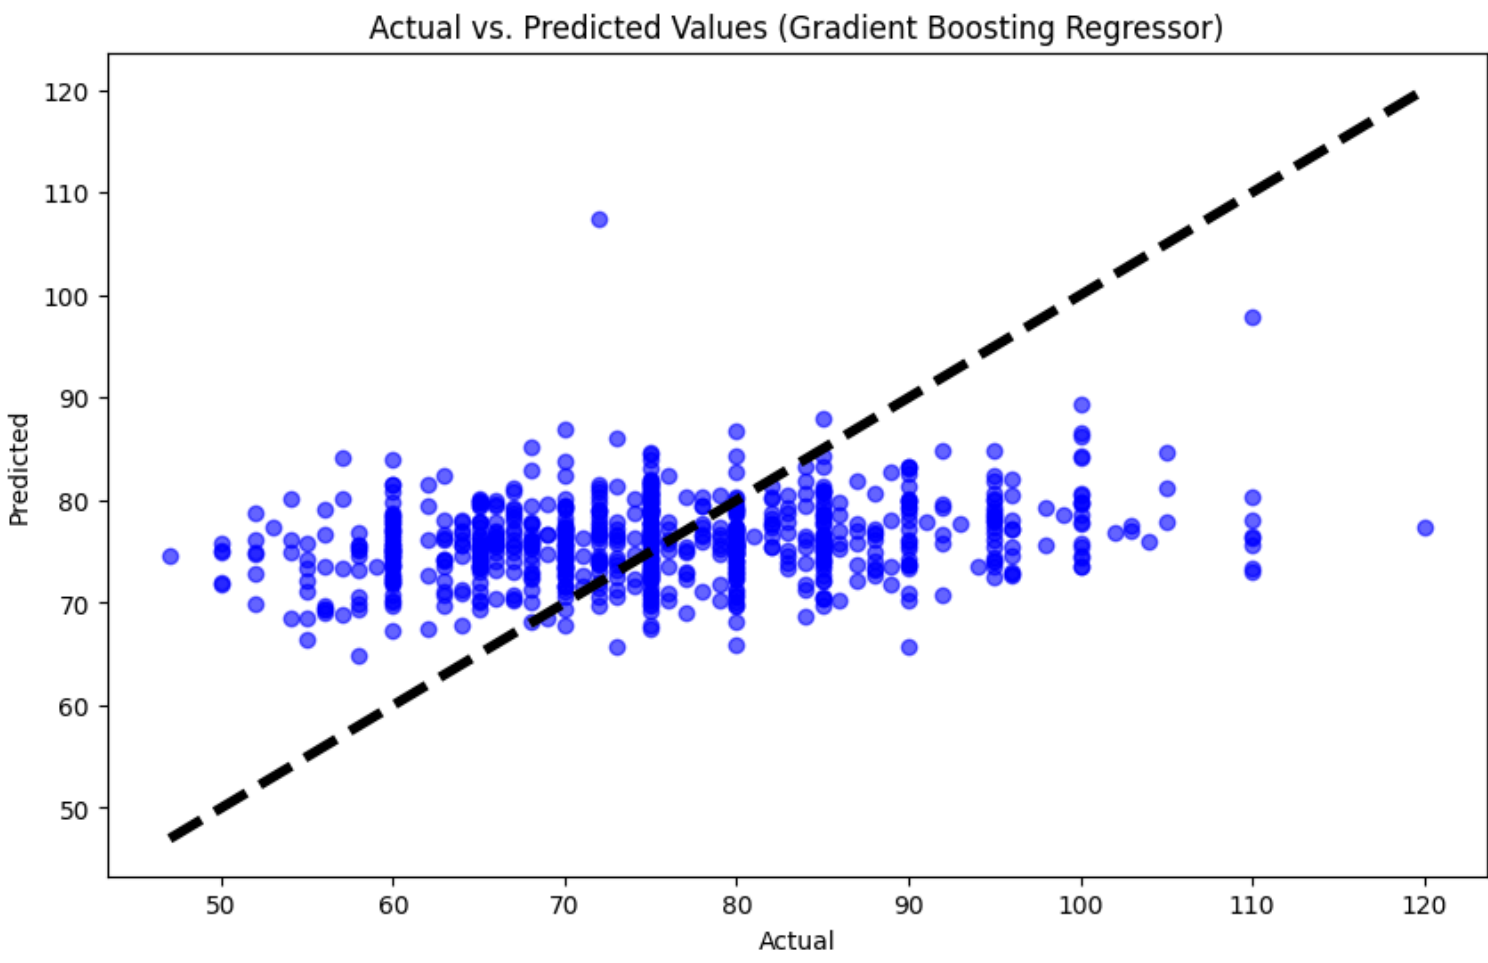
```
plt.figure(figsize=(10, 6))
plt.scatter(Y_test, lin_predictions, alpha=0.6)
plt.plot([Y_test.min(), Y_test.max()], [Y_test.min(), Y_test.max()], 'k--', lw=4) # Diagonal line
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs. Predicted Values( Linear Regression)')
plt.show()
```



```
plt.figure(figsize=(10, 6))
plt.scatter(Y_test, rf_predictions, alpha=0.6, color='green') # Using green for differentiation
plt.plot([Y_test.min(), Y_test.max()], [Y_test.min(), Y_test.max()], 'k--', lw=4) # Diagonal line
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs. Predicted Values (Random Forest Regressor)')
plt.show()
```



```
plt.figure(figsize=(10, 6))
plt.scatter(Y_test, gb_predictions, alpha=0.6, color='blue') # Using blue for differentiation
plt.plot([Y_test.min(), Y_test.max()], [Y_test.min(), Y_test.max()], 'k--', lw=4) # Diagonal line
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs. Predicted Values (Gradient Boosting Regressor)')
plt.show()
```



Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.